

URBANGAON AI REQUIREMENTS DOCUMENT

AI-Powered Internal Tools — Build Brief

Module 1: Daily To-Do & Task Tracker | Module 2: AI Legal Search Tool

Prepared for: AI Tool Development Team

Date: September 01, 2026

Status: Draft v1.0 — For Review

1. Purpose & Overview

This document specifies requirements for two internal AI-powered tools requested by the business: (1) a Daily To-Do / Task Tracker with AI-driven monitoring, and (2) an AI Legal Search Tool for searching across legal documents, judgements, and video content. It captures the original requirements as given, plus additional recommendations to make each tool robust, secure, and genuinely useful in daily operations. The intent is for this to serve as a working brief the AI tool development team can scope, estimate, and build from.

Each module below is broken into: Objective, Core Requirements (as requested), Recommended Enhancements (suggested additions), User Roles, Data & Fields, Workflow, AI Behaviour, and Notifications/Reporting. A shared Technical & Security section and a suggested phased roadmap follow at the end.

2. Module 1 — Daily To-Do List & Task Tracker

2.1 Objective

Give every employee a single place to log their daily/weekly tasks and update status, while an AI layer continuously tracks completion, delays, and time spent — producing manager-ready summaries without anyone filling extra forms.

2.2 Core Requirements (as requested)

- Users can create all of their tasks in the system.
- Users can update and upload the status of each task (e.g. Not Started, In Progress, Completed, Blocked).
- **AI should automatically track:**
 - Which tasks have been completed
 - Which tasks are not completed, along with the reason
 - Time taken to complete each task
- Users can download their Todos date wise

2.3 Recommended Enhancements

The following are suggested additions to make the tool more useful and to reduce manual effort:

- Task creation via natural language — e.g. typing “Follow up with vendor for invoice by Friday” auto-fills title, due date, and priority.
- Recurring tasks (daily/weekly/monthly) so routine work doesn’t need re-entry.
- Sub-tasks and task dependencies (Task B can’t start until Task A is done).
- Effort estimates (in hours), so AI can compare estimated vs. actual time.
- AI-generated daily digest each morning: “You have 3 tasks due today, 1 overdue from yesterday.”
- End-of-day nudge if tasks are still open, prompting the user to update status or give a reason.
- Reason-for-delay categorisation (Waiting on client, Waiting on colleague, Technical issue, Resource unavailable, Deprioritised, Other) so patterns can be analysed across the team, not just read as free text.
- Manager/Team dashboard showing completion rate, average delay, and workload per person — helps identify overloaded team members, not just “who is behind.”
- Attachments and comments on tasks (screenshots, files, links to related work).
- Integration hooks with existing tools if used (email, calendar, Slack/Teams, Jira) so tasks can be created from those without duplicate entry.
- Weekly AI summary report emailed to each user and their manager — completed vs. pending, time-taken trends, recurring blockers.
- Leaderboard / recognition (optional, off by default) for consistent on-time completion — to be used carefully so it doesn't create unhealthy pressure.

2.4 User Roles

Role	Permissions
Employee	Create/edit own tasks, update status, view own history and AI summaries
Manager	View team tasks & dashboards, reassign tasks, view delay-reason analytics
Admin	Full access, manage users/roles, configure categories, export data

2.5 Task Data Fields

Field	Description
Task Title	Short description of the task
Description	Optional detailed notes
Assigned To	Owner of the task (self or team member, if permitted)
Created Date	Auto-captured
Due Date	Set by user
Status	Not Started / In Progress / Completed / Stuck / Cancelled
Reason (if not completed)	Free text + category tag
Start Time / Completion Time	Auto-captured or user-logged, used to compute time taken
Estimated Effort	Optional, in hours
Attachments	Files, links

2.6 AI Behaviour

1. Track status changes and timestamps automatically to compute actual time taken per task.
2. Flag tasks that are overdue or stuck in “In Progress” beyond an expected duration.
3. Read the free-text “reason” field and classify it into a standard category, so trends can be reported (e.g. “38% of delays this month were due to waiting on client inputs”).
4. Generate a daily/weekly plain-language summary per user and per team, instead of requiring anyone to build reports manually.
5. Detect anomalies — e.g. a task marked completed within 2 minutes of creation, or an unusually large number of tasks abandoned by one user — and surface these for manager review rather than silently accepting them.

2.7 Notifications

- Reminder before a task’s due date.
- Alert when a task becomes overdue.
- Daily digest to each user (morning) and daily/weekly summary to managers.
- Optional escalation to manager if a task is overdue by more than a configurable threshold (e.g. 2 days).

3. Module 2 — AI Legal Search Tool

3.1 Objective

Provide a single search interface over the organisation’s legal knowledge base — acts, judgements, contracts, and videos — so users get direct, cited answers instead of manually reading through documents. Any question the AI cannot confidently answer should be routed to an admin so the knowledge base keeps improving.

3.2 Core Requirements (as requested)

- Ability to upload documents such as Acts, Judgements, and YouTube videos (and by extension, other legal reference material).
- Users can search in natural language and get the desired result.
- Questions the AI could not answer should be automatically notified to the admin.

3.3 Recommended Enhancements

The following are suggested additions to strengthen accuracy, trust, and usability:

- Broader content support beyond the requested types: PDFs, scanned documents (with OCR), Word files, web links, audio recordings, and YouTube/other video via transcription.
- Every AI answer includes citations — the specific act/section, judgement name & paragraph, or video timestamp it was drawn from — so users can verify the source rather than trust the answer blindly. This is especially important in a legal context.
- Confidence indication on each answer (e.g. High / Medium / Low confidence, or “no confident match found”) rather than always giving a definitive-sounding answer.
- Document metadata & tagging — jurisdiction, court, year, case type, act name, section number — enabling filtered search (e.g. “only Supreme Court judgements after 2015”).
- Version control for uploaded documents — acts get amended, judgements get overturned/appealed; the system should track supersession so users aren’t shown outdated law as current.
- Bulk upload and folder/category structure for admins adding large volumes of documents at once.
- Search history & saved searches per user, plus the ability to bookmark specific answers or clauses.
- Related-document suggestions (“Users who searched this also viewed...”) to help with legal research beyond the literal query.
- Admin feedback loop: when an unanswered or low-confidence question is resolved by the admin, that answer/document is fed back into the system so the same question is answered correctly next time.
- Audit trail of all searches and answers given, for compliance and quality review — particularly important since legal answers can carry real consequences.
- Access control on sensitive/confidential documents (e.g. internal legal opinions or client-specific contracts) so not all uploaded content is searchable by all users.
- Disclaimer shown on every AI answer that it is a research aid and not a substitute for professional legal advice/verification.

3.4 User Roles

Role	Permissions
Employee/User	Search, view results with citations, save/bookmark, submit questions
Contributor	Upload and tag documents/videos for indexing
Admin	Manage all content, resolve unanswered questions, view analytics, manage access control

3.5 Content Ingestion Workflow

6. Document/video uploaded (by admin or approved contributor) with metadata: type, title, jurisdiction, date, tags.
7. System processes the content — text extraction, OCR (for scans), transcription (for video/audio) — and splits it into indexed, searchable chunks.
8. Content is validated (basic quality/duplicate check) before being marked “live” and searchable.
9. If the document supersedes or amends an existing one, it is linked so users are shown the current version by default, with older versions still accessible for reference.

3.6 Search & Answer Workflow

10. User types a question in natural language (e.g. “What is the notice period required to terminate an employee under the Industrial Disputes Act?”).
11. AI searches indexed content and retrieves the most relevant passages/clauses/video segments.
12. AI generates a direct answer, with citations to the exact source and a confidence indicator.
13. If confidence is low or no relevant content is found, the AI tells the user this plainly and logs the question.
14. Unanswered/low-confidence questions are automatically sent to the admin (via dashboard + email/notification) for review — the admin can then answer directly, upload the missing document, or mark the question as out-of-scope.
15. Once resolved, the admin’s answer/source is added to the knowledge base so future users get it automatically.

3.7 Admin Notification & Resolution

- Real-time or daily-digest notification of unanswered/low-confidence questions, configurable by admin.
- Dashboard listing all pending questions, sorted by frequency (so the most-asked unanswered questions are prioritised).
- Admin can respond, attach/upload the correct source document, or tag a question as resolved/out-of-scope.
- Basic analytics: most-searched topics, most common gaps in the knowledge base, average confidence score over time — helps decide what content to add next.

4. Shared Technical & Security Considerations

4.1 Data Security

- Role-based access control across both tools.
- Encryption of stored documents and data in transit.
- Confidential legal documents restricted to authorised roles only.
- Audit logs for uploads, edits, and admin actions.

4.2 Accuracy & Trust (Legal Search specifically)

- Always cite sources; never present an AI-generated answer as law without a traceable reference.
- Prefer “I don’t have a confident answer” over a guessed answer — false confidence is a bigger risk than an admitted gap in a legal tool.
- Periodic human review of a sample of AI answers for quality assurance.

4.3 Usability

- Both tools should work on desktop and mobile (many employees will log task status or search on the go).
- Simple, minimal-click task entry and search — the tools should reduce effort, not add a new chore.

4.4 Reporting & Analytics

- Task Tracker: completion rate, average delay, workload distribution, delay-reason trends.
- Legal Search: query volume, top topics, unanswered-question backlog, resolution time.

5. Suggested Phased Roadmap

Phase	Scope
Phase 1	Task Tracker: create/update tasks & status, basic AI tracking of completed/pending/reason, simple dashboard. Legal Search: document upload (PDF/text), keyword + AI search with citations, admin notification for unanswered queries.
Phase 2 — Enhance	Task Tracker: recurring tasks, reason categorisation, daily/weekly AI summaries, notifications. Legal Search: video/YouTube transcription support, confidence scoring, version control for documents.
Phase 3 — Scale & Refine	Task Tracker: manager dashboards, anomaly detection, integrations (calendar/Slack/Teams). Legal Search: access control tiers, feedback loop into knowledge base, analytics dashboard.

— End of Document —